










(a) Weak scaling



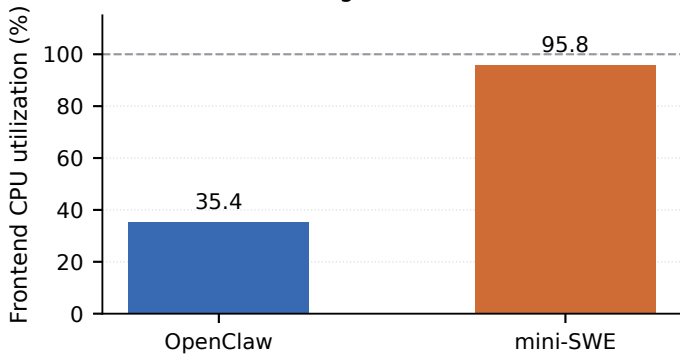
(b) Scaling efficiency



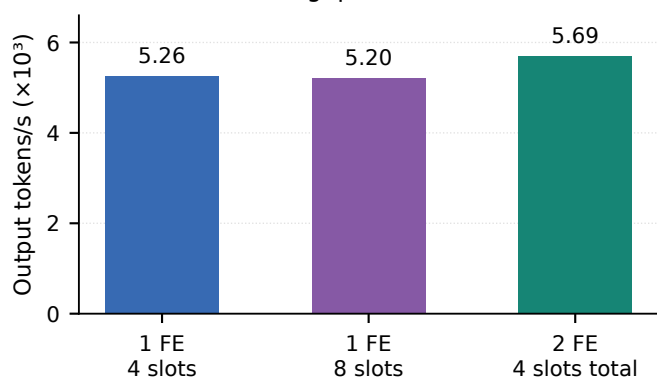
(c) Reuse across scales



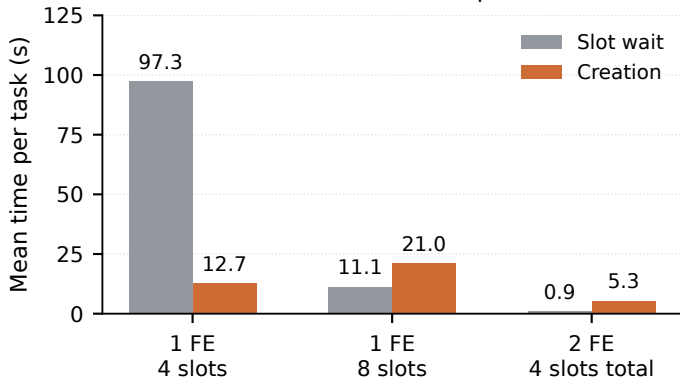
(a) One frontend, eight backends



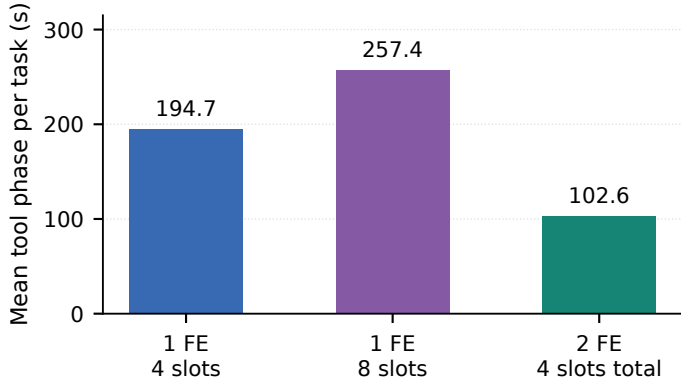
(b) mini-SWE throughput



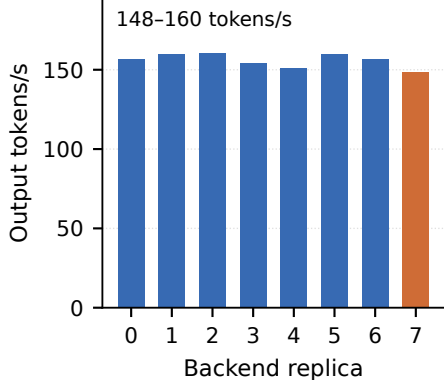
(c) mini-SWE environment setup



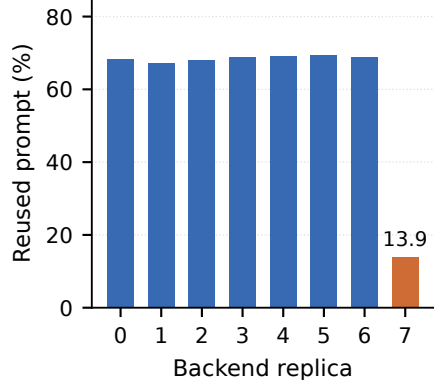
(d) mini-SWE tool execution



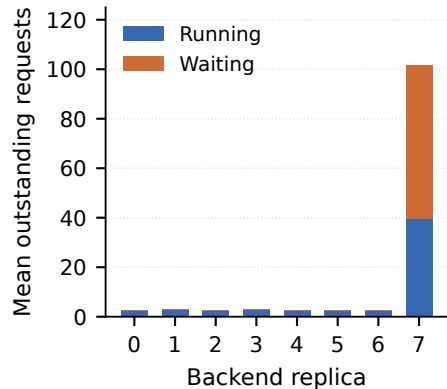
(a) Similar output



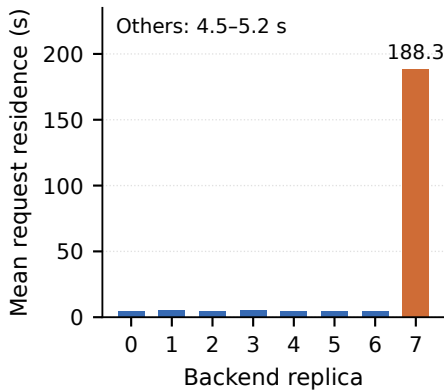
(b) Unequal reuse



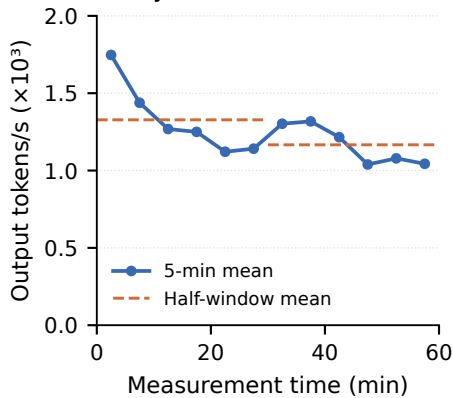
(c) Concentrated backlog



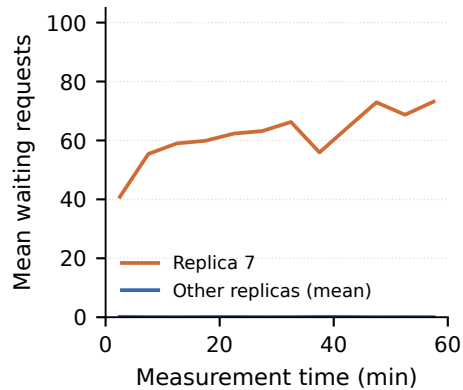
(d) Unequal residence



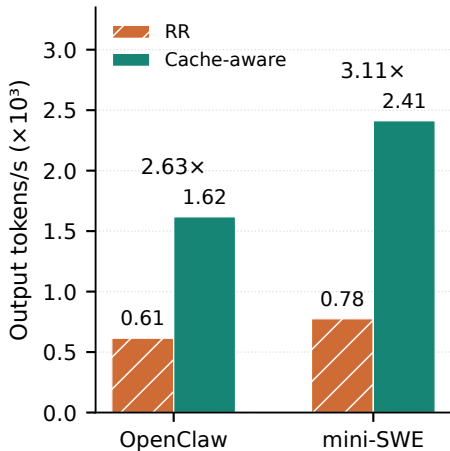
(e) System-wide slowdown



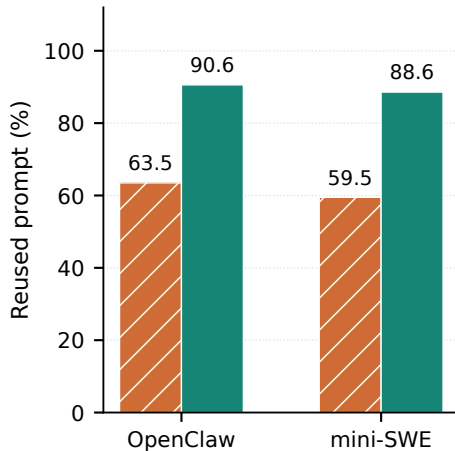
(f) Persistent hotspot



(a) Throughput gain



(b) Prefix reuse



(c) Worst backend queue

