



OpenClaw scaling | Qwen3.6-35B-A3B, TP4, 110-trace pool


(a) Task throughput




(b) Task latency




(c) Inference output throughput




(d) First-token and queue delay



(e) Inference utilization



(f) Prefix-cache reuse



120 s warmup + 1800 s measurement; one repetition per cc. Latencies follow window admissions through natural drain.
web_search uses recorded delay. cc16 is the best observed point in this sweep, not an SLO capacity bound.