



Backend scheduling, cache and output streaming


(a) Scheduler occupancy




(b) Active KV-block occupancy



(c) Engine preemptions



(d) Output streaming slowdown



Free KV blocks can still hold evictable cached prefixes; this gauge does not measure all resident prefix data. Counters span ~1799 s. TPOT is estimated from client streaming chunks, not exact engine per-token timestamps.