



Backend scheduling, cache and output streaming


(a) Scheduler occupancy




(b) KV-cache occupancy



(c) Engine preemptions



(d) Output streaming slowdown



KV occupancy is the reported vLLM gauge, not total allocated GPU memory. Counters span the in-window samples (~1799 s).
TPOT is estimated from client streaming chunks; it is not an exact per-token engine measurement.