



Latency distributions and workload-mix sensitivity

(a) Admission-cohort distributions



(b) Same-trace comparison (n=18)



Boxes: median / IQR, whiskers 1.5 IQR; every observed task is shown. No censored task latency is dropped.
Paired comparison averages repeats per trace, then weights shared traces equally; it does not control cache history or tool outcomes.