



mini-SWE scaling | Qwen3.6-35B-A3B, TP4, 61-trace pool

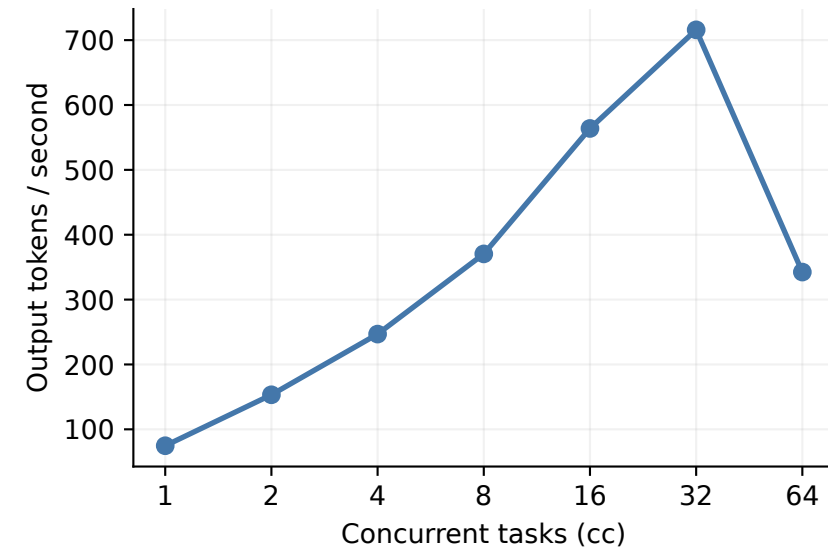
(a) Task throughput



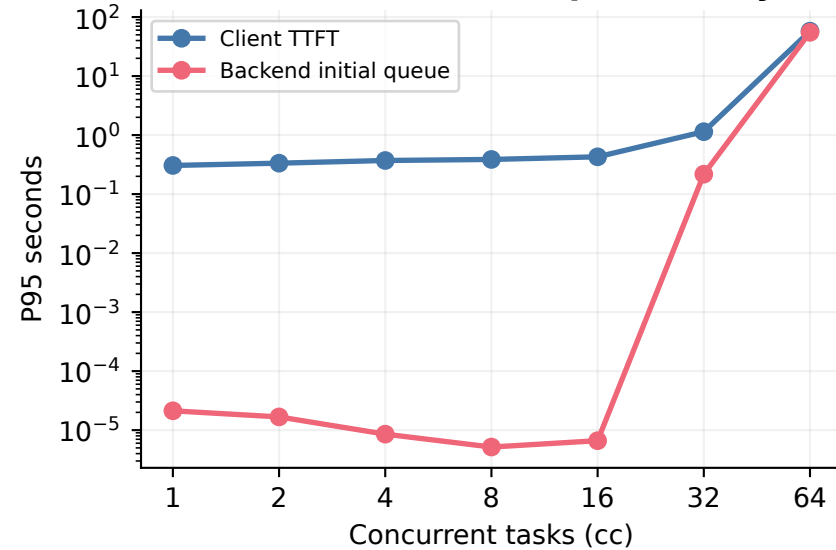
(b) Task latency



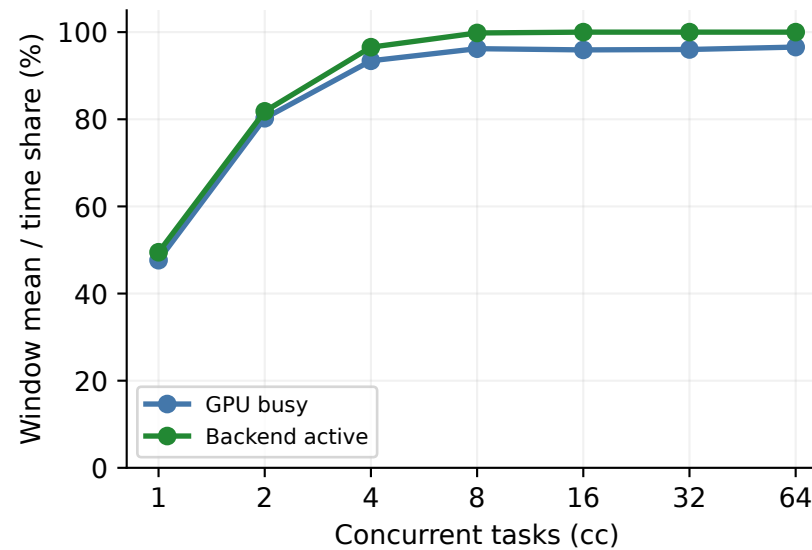
(c) Inference output throughput



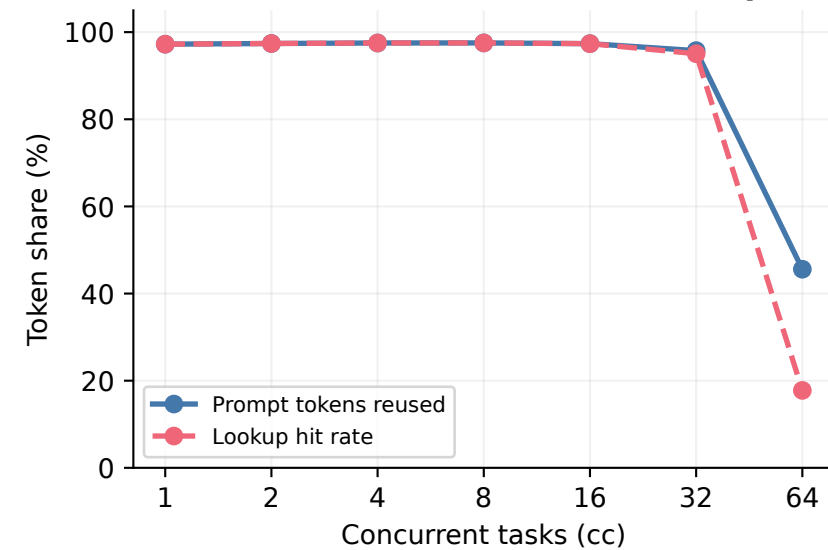
(d) First-token and queue delay



(e) Inference utilization



(f) Cache reuse vs. cache lookups



120 s warmup + 1800 s measurement; one repetition per cc. Latencies follow window admissions through natural drain.
Native bash tools; cc32 has the highest observed throughput. cc64 completions all started in warmup.