



Latency distributions and workload-mix sensitivity

(a) Admission-cohort distributions



(b) First 8 launches (includes warmup)



Boxes: median / IQR, whiskers 1.5 IQR; every observed task is shown. No censored task latency is dropped.
Right: same first 8 tasks, one launch each; mixed warmup/measurement phases. Sensitivity check, not steady-state latency.