


OpenClaw: similar output per backend can hide a concentrated queue



Full-window per-backend means. Prefix reuse is token weighted. Queue and running counts are independent backend samples. RR N4 is nonsteady; RR N8 has two warmup failures and is not declared steady. These are window averages.