





mini-SWE: more frontend resources help; more build slots do not







All measured: 8 TP4 backends, total cc256, sticky routing, 90-minute windows; one run per configuration.
2+8 keeps 4 build slots in total. CPU is averaged per frontend, not summed; tool-error denominator is all window tool calls.

Where frontend time goes: slot waiting, build work, and native tools









Complete measurement-admission cohorts, including natural drain. Boxes: p25-p75; whiskers: p5-p95; diamonds: means. LLM time is task wall time, not exclusive GPU compute. Different closed-loop runs need not have identical task mixtures.

Frontend controls: full-window trajectories




Nonoverlapping 5-minute averages. Slot waiting is grouped by task admission time.
Throughput steadiness does not imply that frontend waiting or task latency has stabilized.

Routing: measured 1+4 comparisons, with the new OpenClaw 1+8 RR window









Hatched: OpenClaw RR N4 nonsteady / N8 warmup failures; mini-SWE RR N4 recovered window with drain timeout.
Blank positions are missing runs, not zero throughput. P2 is excluded. CA = cache-aware; single run per point.

OpenClaw: similar output per backend can hide a concentrated queue







Full-window per-backend means. Prefix reuse is token weighted. Queue and running counts are independent backend samples. RR N4 is nonsteady; RR N8 has two warmup failures and is not declared steady. These are window averages.

OpenClaw RR 4→8 backends: scaling and persistent imbalance



Nonoverlapping 5-minute averages; 60-minute measurements, cc16 per logical worker.
Normalized rates show scale-out efficiency, not recovery to cache-aware/sticky capacity. No steady-state capacity claim.

Frontend/backend scaling: updated measured comparisons



Historical F=B runs use an older HTTP/build policy: capacity references, not controlled frontend-only comparisons.
Mini-SWE 1+4 sticky is missing; no interpolation. New mini-SWE points keep 8 backends and total cc256.